

# Machine Intelligence:: Deep Learning

## Week 2

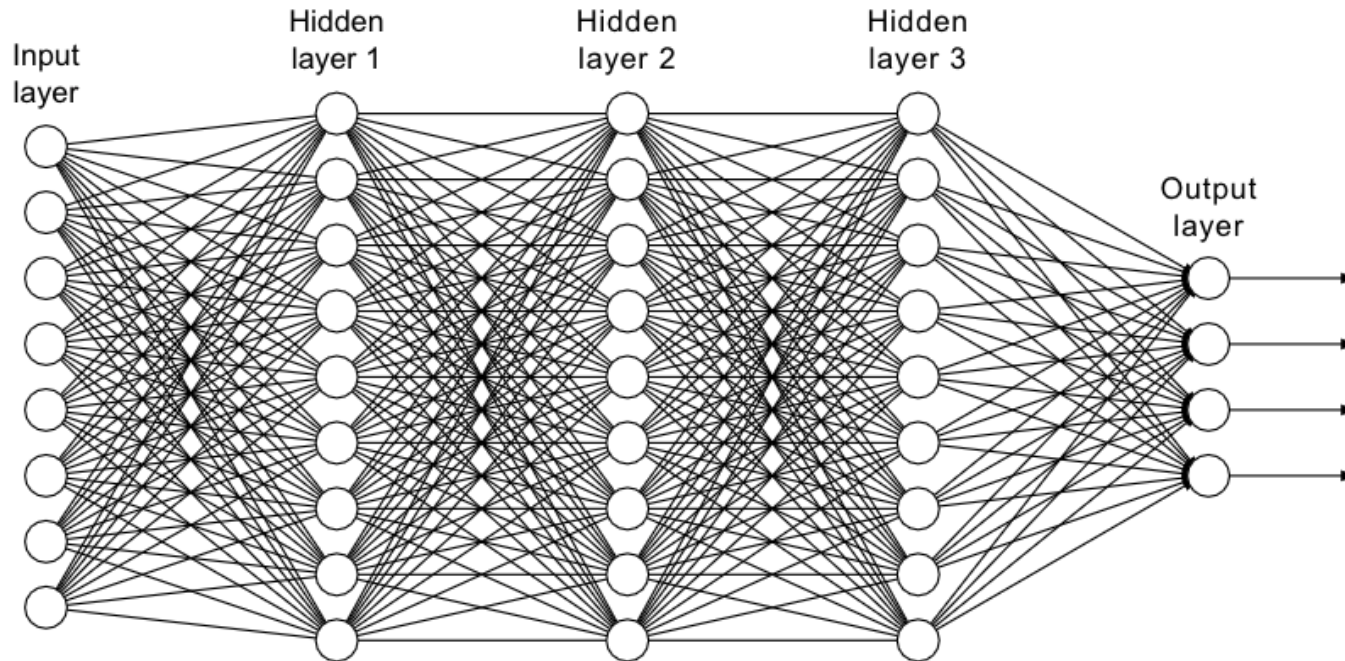
*Beate Sick, Oliver Dürr*

Institut für Datenanalyse und Prozessdesign  
Zürcher Hochschule für Angewandte Wissenschaften

# Topics of today

- A second look on fully connected Neural Networks (fcNN)
- Model development
  - What can we learn from loss curves
    - How to recognize overfitting and underfitting
    - How much data is needed?
    - Does my model learn the right things?
- Convolutional Neural Networks (CNN) for images
  - Motivation for switching from fcNN to CNNs
  - Introduction of convolution
  - ReLu and Maxpooling Layer
  - Biological inspiration of CNNs
  - Building CNNs

# Architecture of a fully connected NN




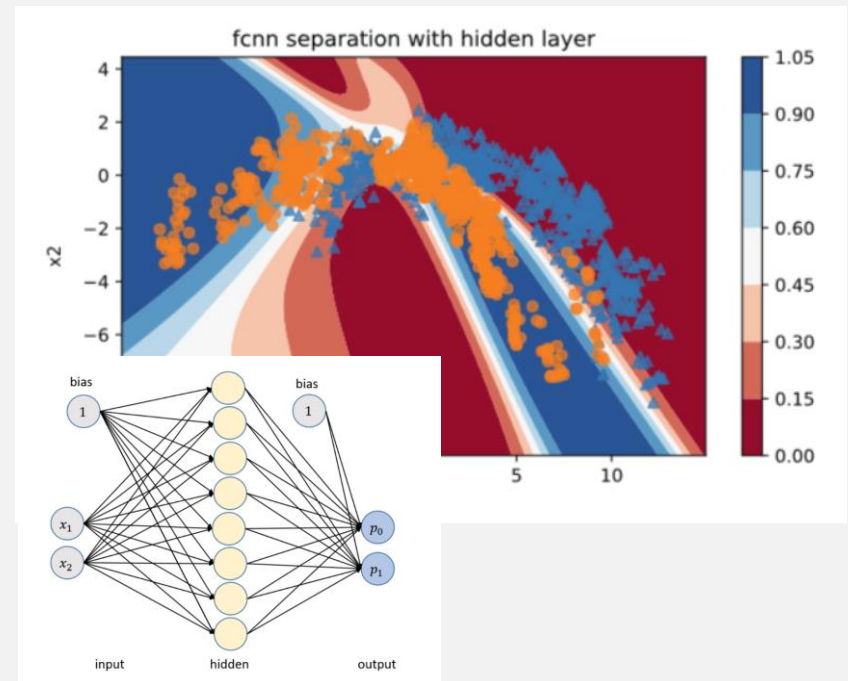


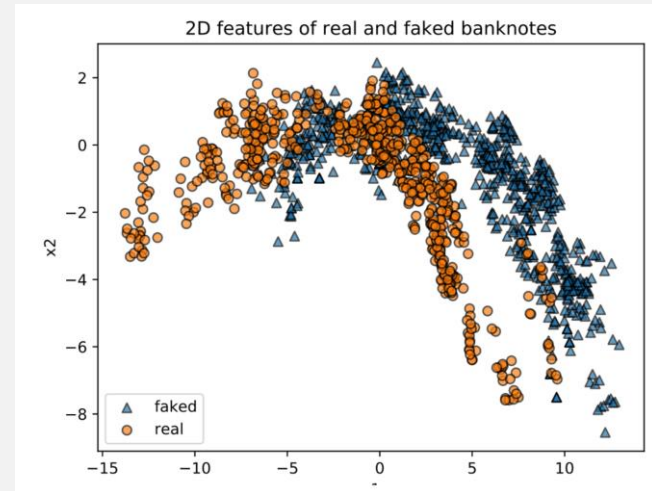
Each neuron in a fcNN gets as input a weighted sum of all neuron activation from one layer below. Different neurons in the same layer have different weights in this weighted sum, which are learned during training.

# Homework


Do exercise NB 02: Classify

banknotes based on 2 features ( $x_1, x_2$ )

[https://github.com/tensorchiefs/dl\\_course\\_2021/blob/master/notebooks/02\\_fcnn\\_with\\_banknote.ipynb](https://github.com/tensorchiefs/dl_course_2021/blob/master/notebooks/02_fcnn_with_banknote.ipynb)



# A deep learning expert at work




# To go deep non-linear activation functions are needed

2 linear layers can be replaced by 1 linear layer -> can't go deep with linear layers!

$$\begin{aligned} z_{21} &= z_{11} \cdot u_{11} + z_{12} \cdot u_{12} = (w_{11} \cdot x_1 + w_{12} \cdot x_2) \cdot u_{11} + (w_{21} \cdot x_1 + w_{22} \cdot x_2) \cdot u_{12} \\ &= x_1 \cdot (w_{11} \cdot u_{11} + w_{21} \cdot u_{12}) + x_2 \cdot (w_{12} \cdot u_{11} + w_{22} \cdot u_{12}) \end{aligned}$$


$$z_{11} = w_{11} \cdot x_1 + w_{12} \cdot x_2$$

$$z_{12} = w_{21} \cdot x_1 + w_{22} \cdot x_2$$



=

$$z_{21} = v_{11} \cdot x_1 + v_{12} \cdot x_2$$



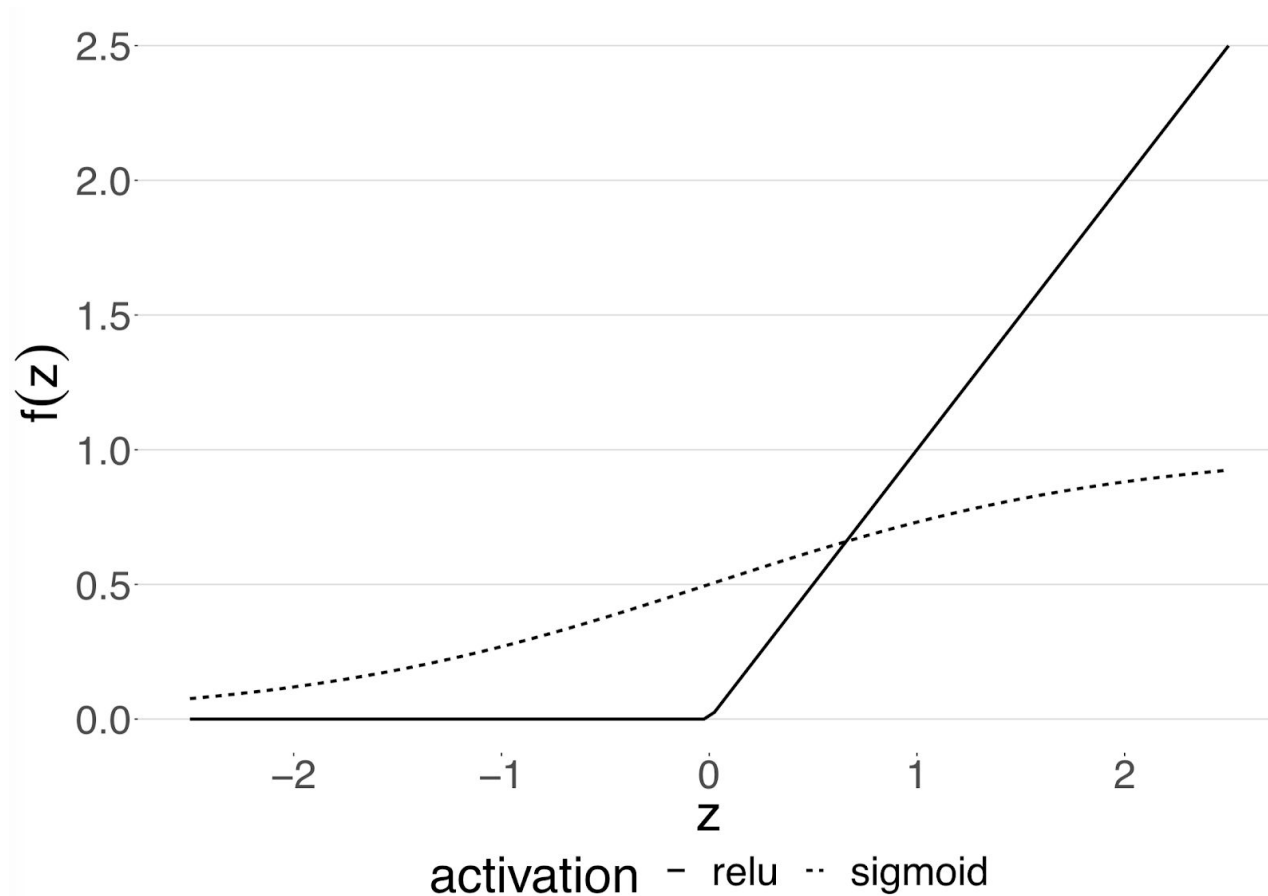
$$v_{11} = w_{11} \cdot u_{11} + w_{21} \cdot u_{12}$$

$$v_{12} = w_{12} \cdot u_{11} + w_{22} \cdot u_{12}$$

$$v_{21} = w_{11} \cdot u_{21} + w_{21} \cdot u_{22}$$

$$v_{22} = w_{12} \cdot u_{21} + w_{22} \cdot u_{22}$$

# Comon non-linear activation function



The sigmoid has small gradients for values far away from zero.  
ReLU clips values below zero and let values  $> 0$  pass unchanged.


# Model development

## Overfitting and underfitting



# "Perfect" Vs. "Simple" classifier

"Perfect"  
classifier



"Simple"  
classifier

Which is better?

Check on a test-set (cross validation).

# Cross validation of the “simple” classifier

Training set:


6/29=20%

misclassification

Test set:

2/25=8%


misclassification



# Cross validation of the “Perfect” classifier

Training set:  
0% misclassification

Test set:  
8/25=24%  
misclassification




# Types of Errors

- **Training error or naïve error or in-sample-error:**
  - Error on data that have been used to train the model
- **Generalization error or test error or out-of-sample-error:**
  - Error on previously unseen records (out of sample)

## **Over-fitting phenomenon:**


- Model fits the training data well (small training error) but shows high generalization error

# What is the right level of complexity



Remark: In DL the models are often very flexible and show overfitting when trained over many epochs – early stopping or regularization are needed.

# Best practice: Split in Train, Validation, and Test Set




Best practice: Lock an extra **test data set** away, and use it only at the very end, to evaluate the chosen model, that performed best on your validation set.

Reason: **When trying many models, you probably overfit on the validation set.**

Determine performance metrics, such as MSE, to evaluate the predictions **on new validation or test data**

# What can loss curves tell us?


Very common check: Plot loss in train and validation data vs epoch of training.



- If training loss does not go down to zero: model is not flexible enough
- In case of overfitting (validation loss  $\gg$  train loss): regularize model

# What can loss curves tell us?

Less common check: Plot loss in train and validation data vs amount of train data.



possible cures for overfitting:

- more data (if curves are still approaching)
- more representative data (if gap is constant)
- less complex model (regularization)
- bagging (and boosting)

possible cures for underfitting:



- Use more complex model (less regularization)
- Use other model structure
- boosting



# Model training process relies on “the big lie”

$P(\text{train}) = P(\text{test}) = P(\text{test})$ : The “big lie”

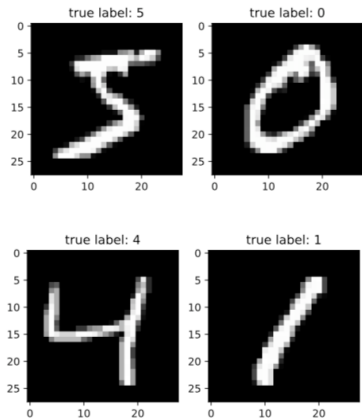
- In traditional machine learning and DL we usually assume that the train and test data do come from the same distribution.
- If this assumption is not true, a seemingly high performant model breaks down and cannot see the elephant in the room.
- We need models that can flag uncertain predictions or even better learn relevant (causal) features that do not depend on distribution.




Remark: we will discuss this in more detail in later lectures and also learn about some cures.

Fully connected NN for image data  
Why not?

# A fcNN for MNIST data




The first four digits of the MNIST data set - each image consisting of  $28 \times 28 = 784$  pixels



A fully connected NN with 2 hidden layers.

For the MNIST example, the input layer has 784 values for the  $28 \times 28$  pixels and the output layer has 10 nodes for the 10 classes.

# What is going on in a 1 layer fully connected NN?

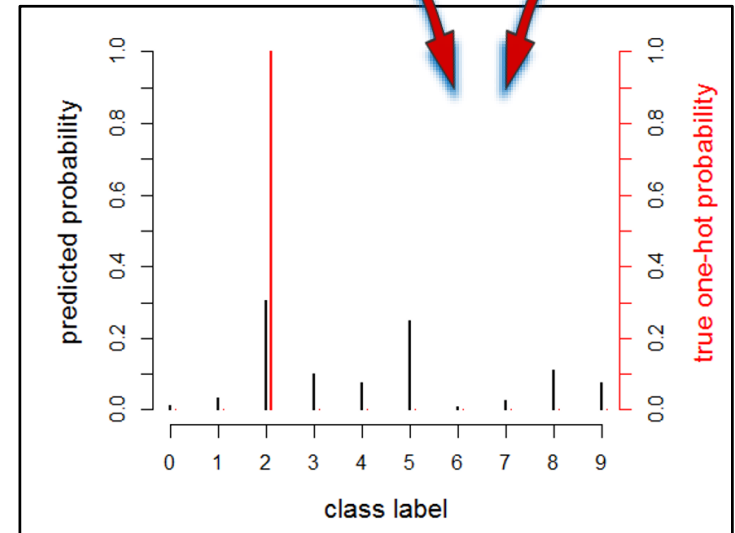


Cost C or Loss = cross-entropy averaged over all images in mini-batch


$$C = \frac{1}{N} \sum_i D(\mathbf{p}_i, \mathbf{y}_i)$$

Cross-Entropy

$$D(\mathbf{p}, \mathbf{y}) = - \sum_{k=1}^{10} y_k \cdot \log(p_k)$$



# Exercise: Does shuffling disturb a fcNN?



Use fcNN for MNIST:

[https://github.com/tensorchiefs/dl\\_course\\_2020/blob/master/notebooks/03\\_fcnn\\_mnist.ipynb](https://github.com/tensorchiefs/dl_course_2020/blob/master/notebooks/03_fcnn_mnist.ipynb)

Investigate if shuffling disturbs the fcNN for MNIST:


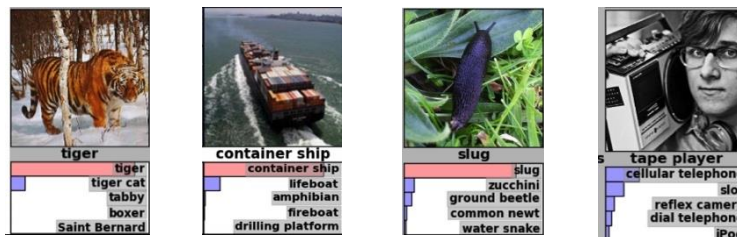
[https://github.com/tensorchiefs/dl\\_course\\_2020/blob/master/notebooks/04\\_fcnn\\_mnist\\_shuffled.ipynb](https://github.com/tensorchiefs/dl_course_2020/blob/master/notebooks/04_fcnn_mnist_shuffled.ipynb)

# Convolutional Neural Networks

## SoA for image data

# Recall: Imagenet challenge

1000 classes  
1 Mio samples



Human: 5% misclassification

Only one non-CNN approach in 2013

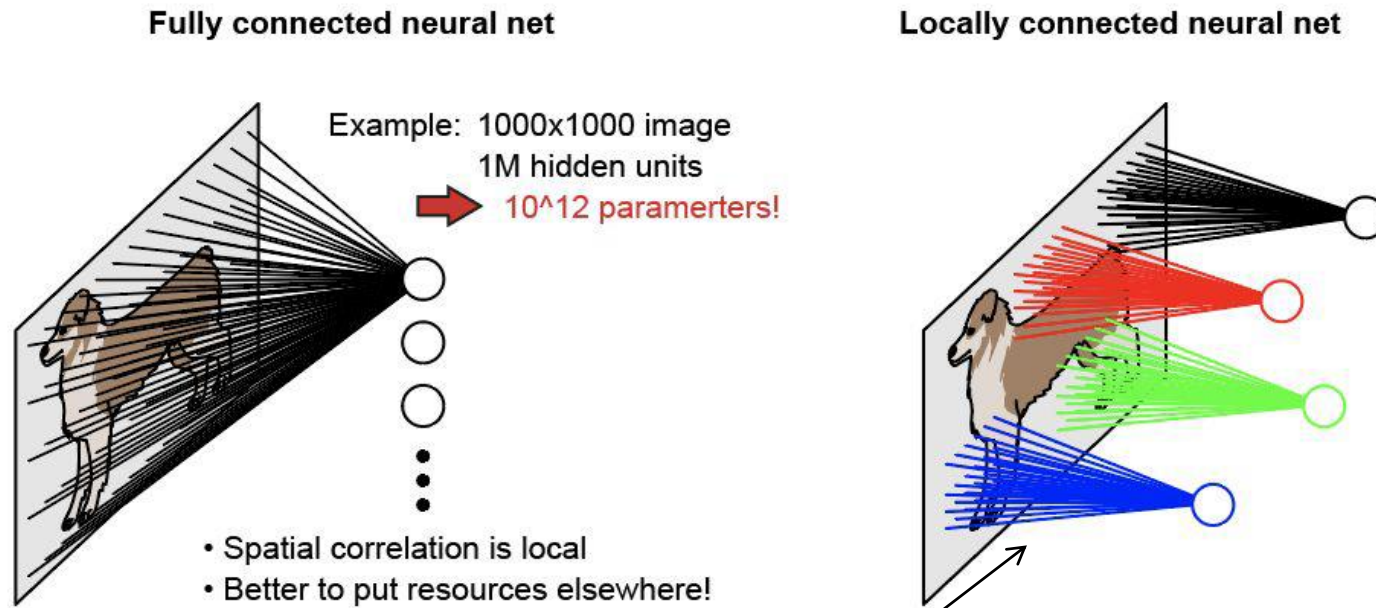
GoogLeNet 6.7%

A. Krizhevsky  
first CNN in 2012  
**Und es hat zoom gemacht**

2015: It gets tougher

4.95% Microsoft ([Feb 6](#) surpassing human performance 5.1%)  
4.8% Google ([Feb 11](#)) -> further improved to 3.6 (Dec)?  
4.58% Baidu (May 11 [banned due too many submissions](#))  
3.57% Microsoft (Resnet winner 2015)

# Convolution extracts local information using few weights

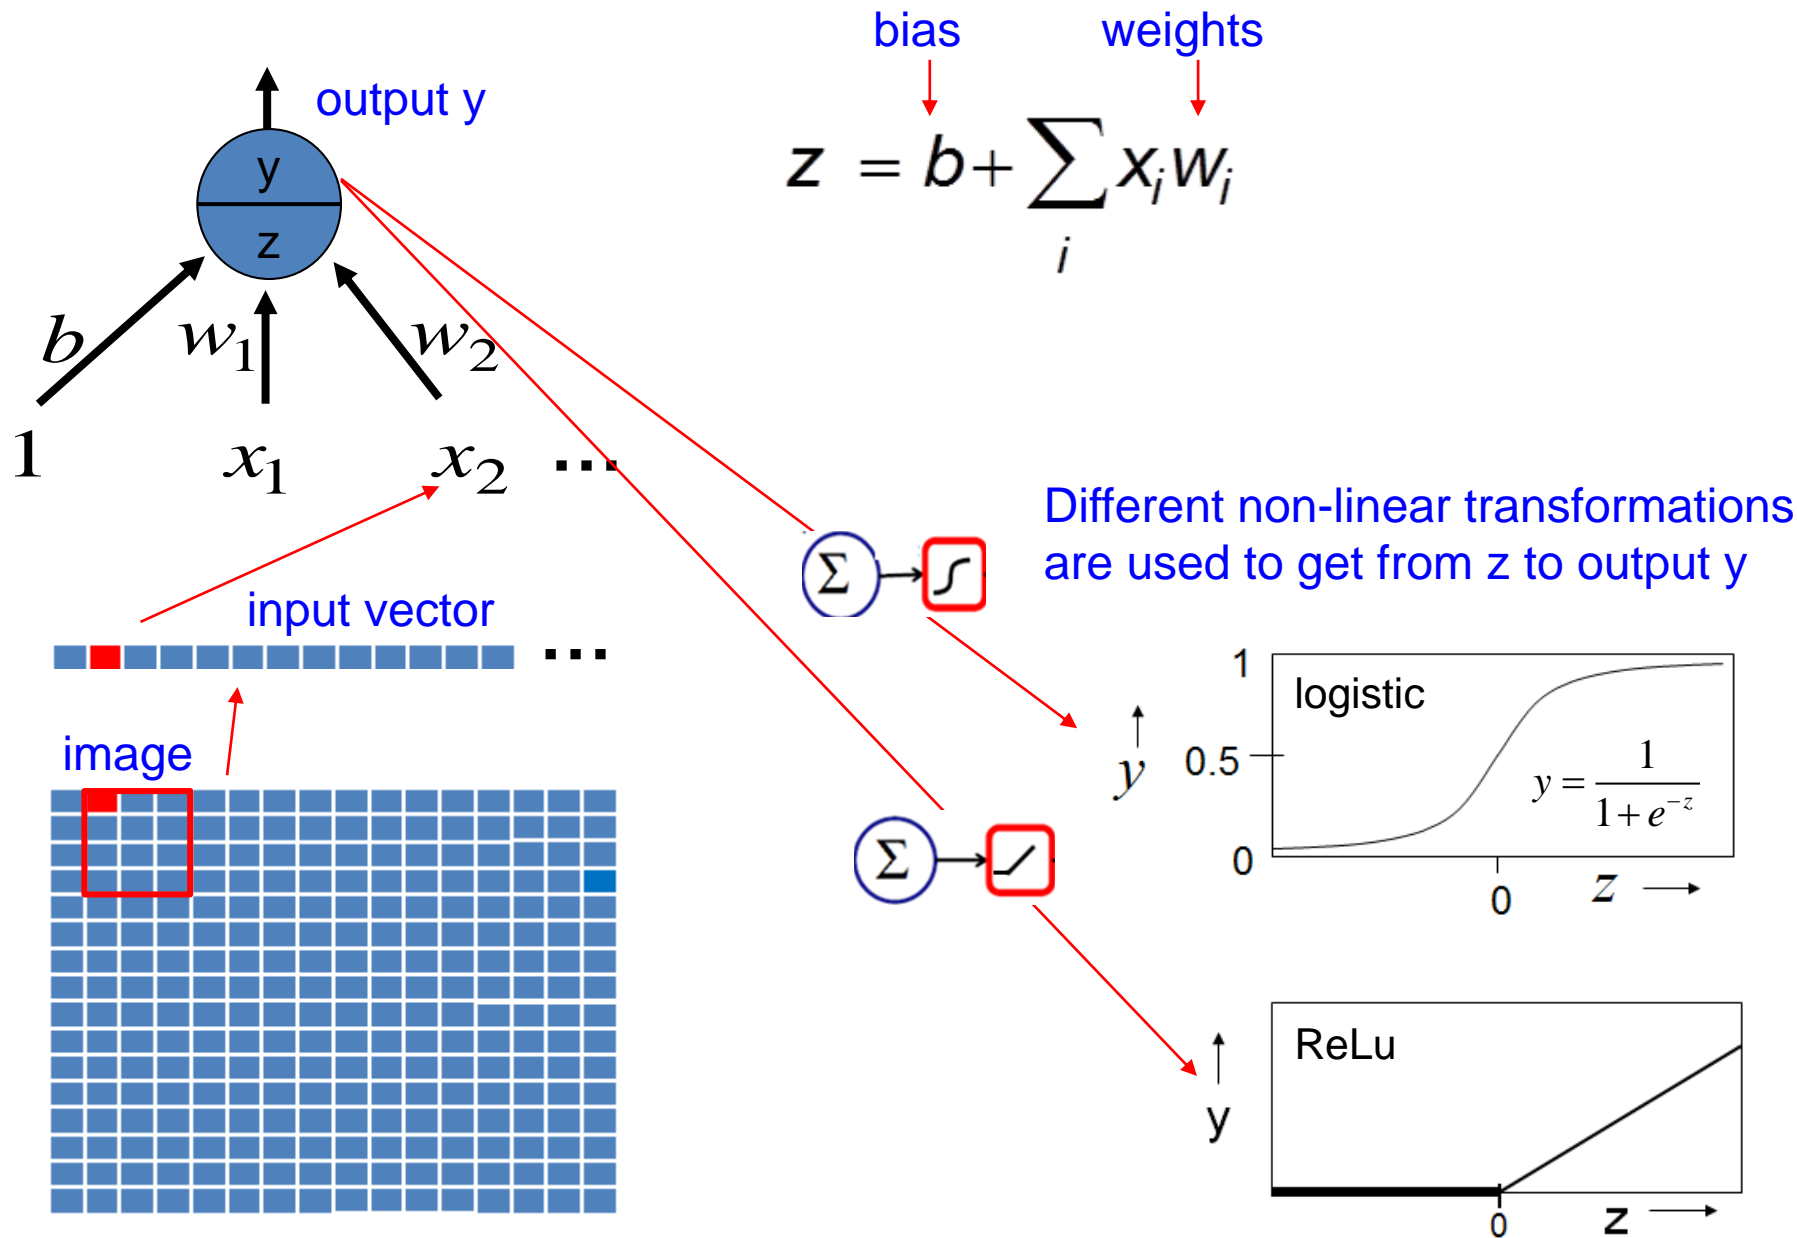


## Shared weights:

by using the **same weights for each patch** of the image we need much **less parameters** than in the fully connected NN and get from each patch the same kind of **local feature information** such as the presence of a edge.



# An artificial neuron





# Convolutional networks use neighborhood information and replicated local feature extraction

In a locally connected network the calculation rule

$$z = b + \sum_i x_i w_i$$



Pixel values in a small image patch are element-wise multiplied with weights of a small filter/kernel:




The filter is applied at each position of the image and it can be shown that the **result is maximal if the image pattern corresponds to the weight pattern.**

The results form again an image called **feature map (=activation map)** which shows at which position the feature is present.

feature/activation map



# Applying the same 3x3 kernel at each image position






Applying the 3x3 kernel on a certain position of the image yields one pixel within the activation map where the position corresponds to the center of the image patch on which the kernel is applied.

## Convolutional networks use neighborhood information and replicated local feature extraction

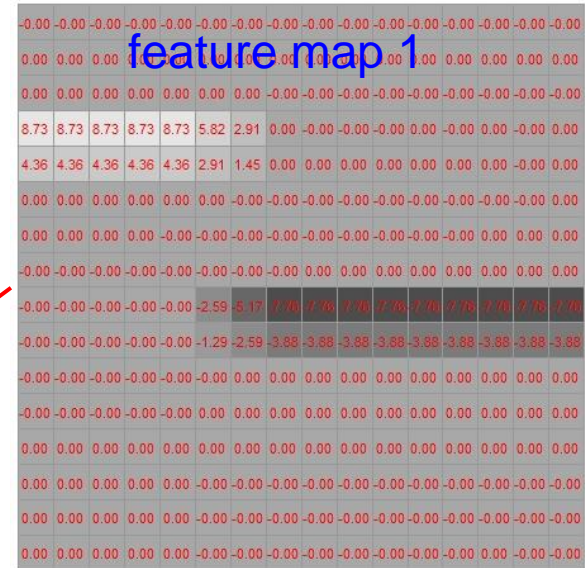
filtering =  
convolution

kernel 1

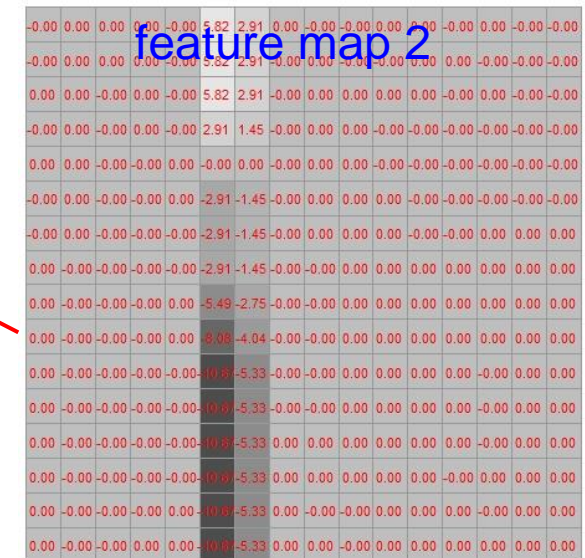
image



feature map 1



feature map 2



The weights of each filter are randomly initiated and then adapted during the training.


## Exercise: Do one convolution step by hand



The kernel is 3x3 and is applied at each valid position  
– how large is the resulting activation map?

The small numbers in the shaded region are the kernel weights.  
Determine the position and the value within the resulting activation map.

3	3	2	1	0
$0_0$	$0_1$	$1_2$	3	1
$3_2$	$1_2$	$2_0$	2	3
$2_0$	$0_1$	$0_2$	2	2
2	0	0	0	1






# Example of designed Kernel / Filter




$$\begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix}$$

Horizontal Sobel kernel




Applying a vertical edge detector kernel

# CNN Ingredient I: Convolution




Zero-padding to achieve  
**same** size of feature and input



no padding to only use  
**valid** input information

The *same* weights are used at each position of the input image.

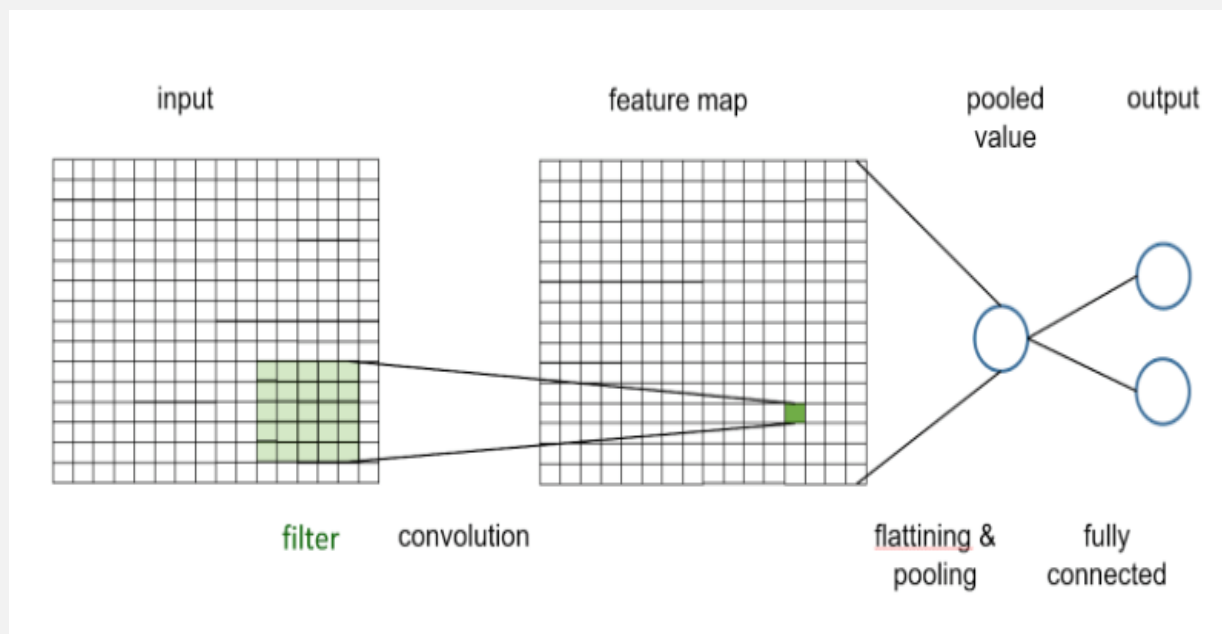
# Building a very simple CNN with keras



```
model = Sequential()
model.add(Convolution2D(1, (5,5), # one 5x5 kernel
                        padding='same', # zero-padding to preserve size
                        input_shape=(pixel,pixel,1)))
model.add(Activation('linear'))
# take the max over all values in the activation map
model.add(MaxPooling2D(pool_size=(pixel,pixel)))
model.add(Flatten())
model.add(Dense(2))
model.add(Activation('softmax'))
```




# Exercise: Artstyle Lover



Open NB in: [https://github.com/tensorchiefs/dl\\_course\\_2020/blob/master/notebooks/05\\_cnn\\_edge\\_lover.ipynb](https://github.com/tensorchiefs/dl_course_2020/blob/master/notebooks/05_cnn_edge_lover.ipynb)


## Convolution layer with a 1-channel input and 6 kernels




Convolution of the input image with 6 different kernels results in 6 activation maps.  
If the input image has only one channel, then each kernel has also only one channel.

# Animated convolution with 3 input channels

3 color channel input image

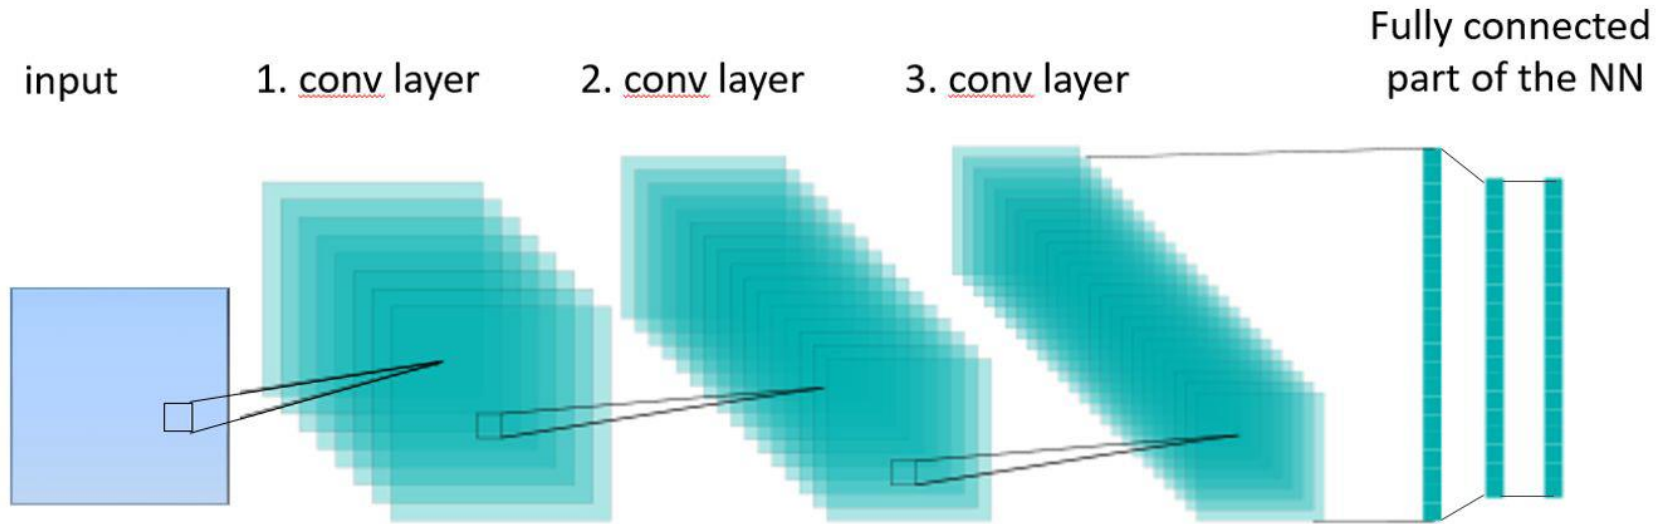


## Convolution layer with a 3-channel input and 6 kernels


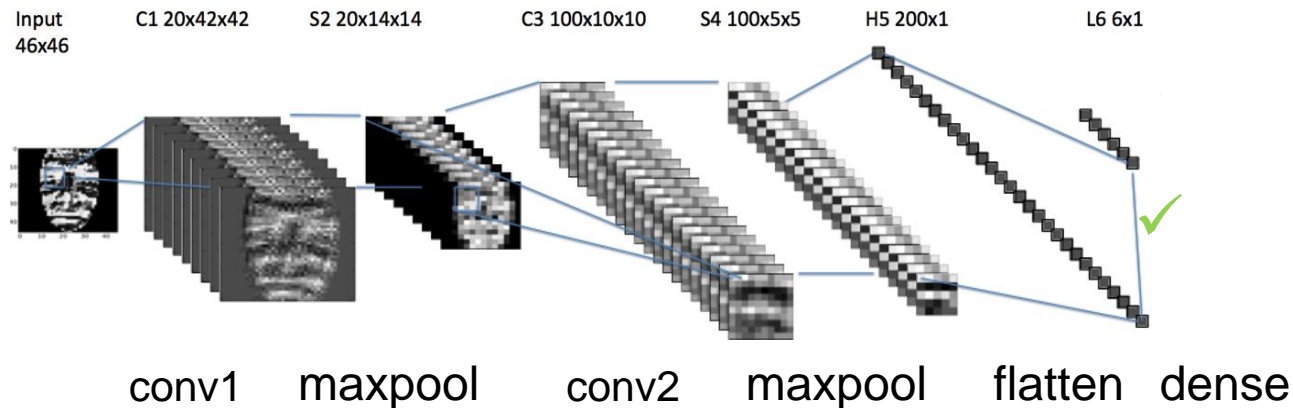


Convolution of the input image with 6 different kernels results in 6 activation maps. If the input image has 3 channels, then each filter has also 3 channels.

# A CNN with 3 convolution layers




# CNN ingredient II: Maxpooling Building Blocks reduce size



Simply join e.g. 2x2 adjacent pixels in one by taking the max.  
→ less weights in model  
→ Less train data needed  
→ increased performance

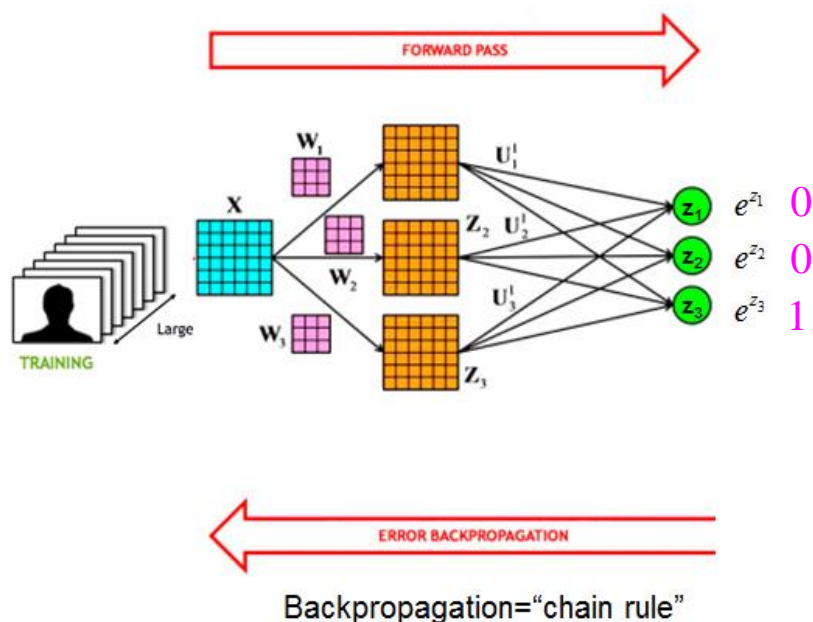
Hinton: „The pooling operation used in convolutional neural networks is a big mistake and the fact that it works so well is a disaster“

# Typical shape of a classical CNN



Spatial resolution is decreased e.g. via max-pooling while more abstract image features are detected in deeper layers.

# Training of a CNN is based on gradient backpropagation



For the training we need the **observed label** for each image which we then compare with the **output** of the CNN.

We want to **adjust the weights** in a way so that difference between true label and output is minimal.

Minimize Loss-function:


**L=distance(observed, output(w))**

Learning is done by weight updating:

$$w_i^{(t)} = w_i^{(t-1)} - \underset{\substack{\uparrow \\ \text{learning rate}}}{l^{(t)}} \frac{\partial L(w)}{\partial w_i} \bigg|_{w_i = w_i^{(t-1)}}$$



# Exercise: Does shuffling disturb a CNN?




Open NB in: [https://github.com/tensorchiefs/dl\\_course\\_2020/blob/master/notebooks/06\\_cnn\\_mnist\\_shuffled.ipynb](https://github.com/tensorchiefs/dl_course_2020/blob/master/notebooks/06_cnn_mnist_shuffled.ipynb)

# fcNN versus CNNs – some aspects

- A fcNN is good for tabular data, CNNs are good for ordered data (eg images).
- In a fcNN the order of the input does not matter, in CNN shuffling matters.
- A fcNN has no model bias, a CNN has the model bias that neighborhood matters.
- A node in one layer of a fcNN corresponds to one feature map in a convolution layer:
- In each layer of a fcNN connecting  $p$  to  $q$  nodes, we learn  $q$  linear combinations of the incoming  $p$  signals, in each layer of a CNN connecting  $p$  channels with  $q$  channels we learn  $q$  filters (each having  $p$  channels) yielding  $q$  feature maps

# Homework: Develop a CNN for cifar10 data



Develop a CNN to classify cifar10 images (we have 10 classes)

Investigate the impact of standardizing the data on the performance

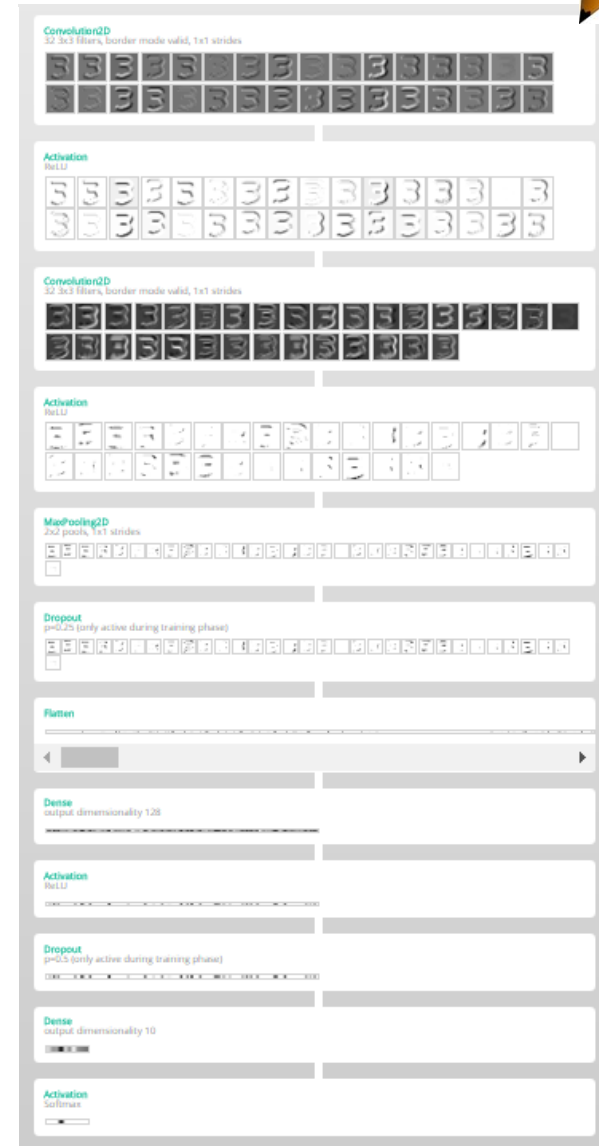

Notebook for homework will soon be on the course webpage.

# Homework: a simple live CNN: fill the gaps (optional)

Follow the first link ["live cnn in browser"](#)

<https://transcranial.github.io/keras-js/#/mnist-cnn>

```
model = Sequential()  
model.add(Convolution2D(... , ... , ... ,  
                        border_mode='valid',  
                        input_shape=(... , ... ),  
                        dim_ordering='tf'))  
model.add(Activation('...'))  
model.add(Convolution2D(... , ... , ... ,  
                        border_mode='valid',  
                        dim_ordering='tf'))  
model.add(Activation('...'))  
model.add(MaxPooling2D(pool_size=(... , ...),  
                        border_mode='valid',  
                        dim_ordering='tf'))  
model.add(Dropout(...))  
model.add(Flatten())  
model.add(Dense(...))  
model.add(Activation('relu'))  
model.add(Dropout(...))  
model.add(Dense(...))  
model.add(Activation('softmax'))
```



# Summary

- Use loss curves to detect overfitting or underfitting problems
- NNs work best when respecting the underlying structure of the data.
  - Use fully connected NN for tabular data
  - Use convolutional NN for data with local order such as images
- CNNs exploit the local structure of images by local connections and shared weight (same kernel is applied at each position of the image).
- Use the relu activation function for hidden layers in CNNs.